

Bicycle project crowd evaluation

Hugo Dalbousiere

July 6, 2021

Abstract

In this experiment we will analyze the accuracy of a crowd of annotators which had to answer the following binary question: “Do you see a bicycle?”.

1 Introduction

Each annotator is given a set of tasks (work package) that it has to solve. Every task is pointing to a specific image hosted on Amazon S3. Based on this image, the annotator should determine if it contains a bicycle or not. After the annotator has handled all its assigned tasks, the results are sent back to the main server and aggregated for later data analysis. The objective of this experiment will be to analyse the results of the annotators which are stored in a JSON file.

2 JSON file format inspection

It is essential to be comfortable with the format of the given data set in order to exploit them efficiently and accurately. The JSON file weights about 64MB and is not prettified, making it challenging to visualize it in a proper way. After prettifying the file and stripping off unnecessary data, we have a sample file that is adequate to start taking a look at the different properties of the file (see figure 2).

```
1 {
2   "results": {
3     "root_node": {
4       "gui_type": "discrete_answer",
5       "results": {
6         "7e8984b6-dff7-4015-865a-b721a2faf681": {
7           "results": [
8             {
9               "task_input": {
10                 "image_url": "https://qm-auto-annotator.s3.eu-central-1.
amazonaws.com/bicycles/img_4686.jpg"
11               },
12               "created_at": "2021-02-25T14:08:11.319438+00:00",
13               "workpackage_total_size": 5,
14               "loss": 0.0,
15               "project_node_input_id": "7e8984b6-dff7-4015-865a-
b721a2faf681",
16               "project_node_output_id": "0000439a-96ac-4bd4-8753-
a4baa229ecf2",
17               "task_output": {
18                 "answer": "no",
19                 "cant_solve": false,
20                 "corrupt_data": false,
21                 "duration_ms": 997
22               },
23               "user": {
24                 "vendor_id": "vendor_01",
25                 "id": "08af8775-a72c-4c59-b60f-9ce7df04fa92",
26                 "vendor_user_id": "annotator_12"
27               },
28               "root_input": {
29                 "image_url": "https://qm-auto-annotator.s3.eu-central-1.
amazonaws.com/bicycles/img_4686.jpg"
30               },
31               "project_root_node_input_id": "7e8984b6-dff7-4015-865a-
b721a2faf681"
32             }
33           ],
34           "gui_type": "discrete_answer"
35         }
36       }
37     }
38   }
39 }
```

Figure 2: Content of a prettified sample of the JSON file

The file contains two first higher levels that not bringing much value: “results” and “root_node”. However, the second “results” property contained in the “root_node” property is a key element. It contains a set of key and values: the key is the unique identifier of a work package and the value contains an array of results. Each entry in this array is the result of a single task. Each task contains some metadata, such as the annotator id (user id), crowd (vendor), task duration and the image url, but most importantly it contains the answer returned by the annotator: “yes” or “no”.

3 Data Preparation

A JSON file can be manipulated using Python thanks to the “json” library, but it is easier to manipulate data models rather than dictionaries as they have no predefined structure. Data models¹ can be assimilated to classes in Object Oriented Programming (OOP) and have many advantages (re-usability, maintainability, code readability...). For this use case, one of the most interesting benefits of data models is having a loose coupling to the format of the JSON file: if the data format of the file were to change, the only adaptation needed would be how the data models are created, but the business logic leveraging these data models would remain unchanged. Two data models will be created: Annotator and Task (see figure 3)

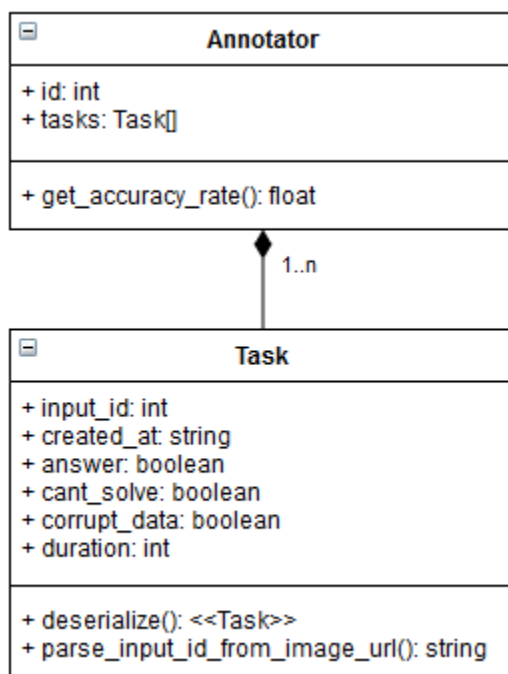


Figure 3: UML diagram of the data models

¹See “Data model” from the Python documentation: <https://docs.python.org/3/reference/datamodel.html>

4 Analysis

The data set contains 22 annotators. The minimum annotation time is -99999ms, while the maximum annotation time is 42398ms. These numbers alone do not give much insight on the whole data set, and having a negative time is most likely the outcome of an error during the aggregation of the results. The histogram in figure 4 allows us to have an overview of the distribution of the duration time.

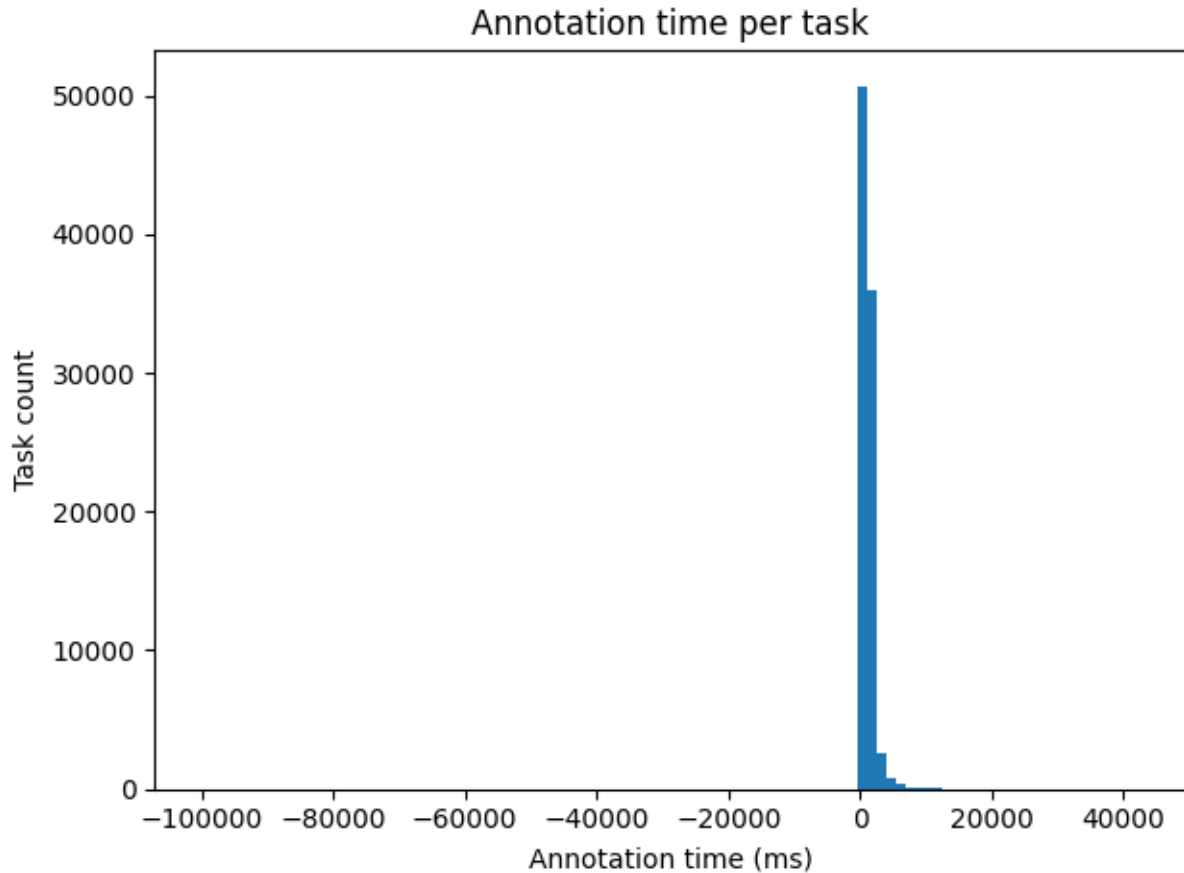


Figure 4: Annotation time per task

While it is clear that most of the tasks take between 0 and 5000ms to execute, the histogram is hard to read because of the very few outlier values in the data set. In the next histogram (see figure 5), outlier values (< 0 ms or > 5000 ms) are excluded.

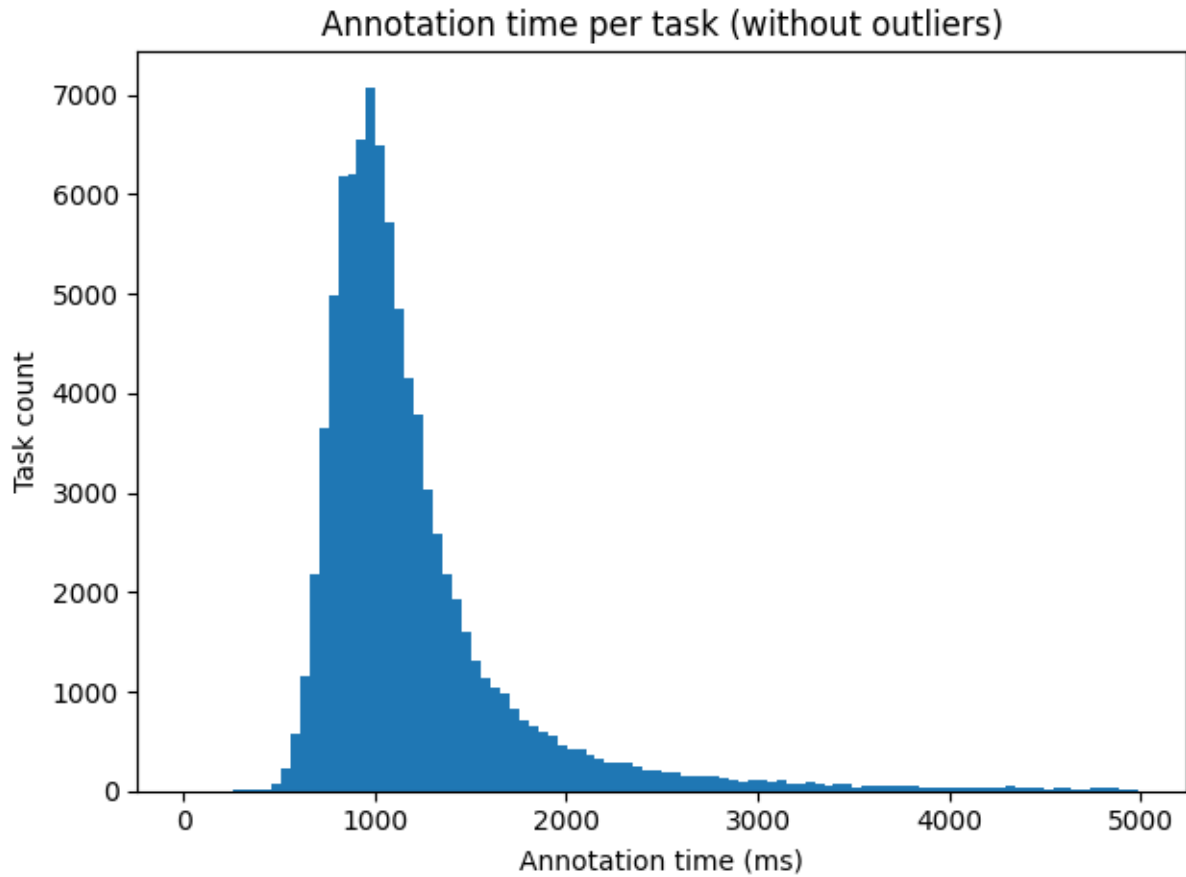


Figure 5: Annotation time per task (without outliers)

It's now clear that most of the tasks take between 500 and 1500ms to complete. The average annotation time is 1284ms, but it is more accurate to use the median time as we have outlier values. The median average annotation time is 1058ms.

Annotators each produced between 170 to 7596 results (see figure 6). A result corresponds to an annotated task.

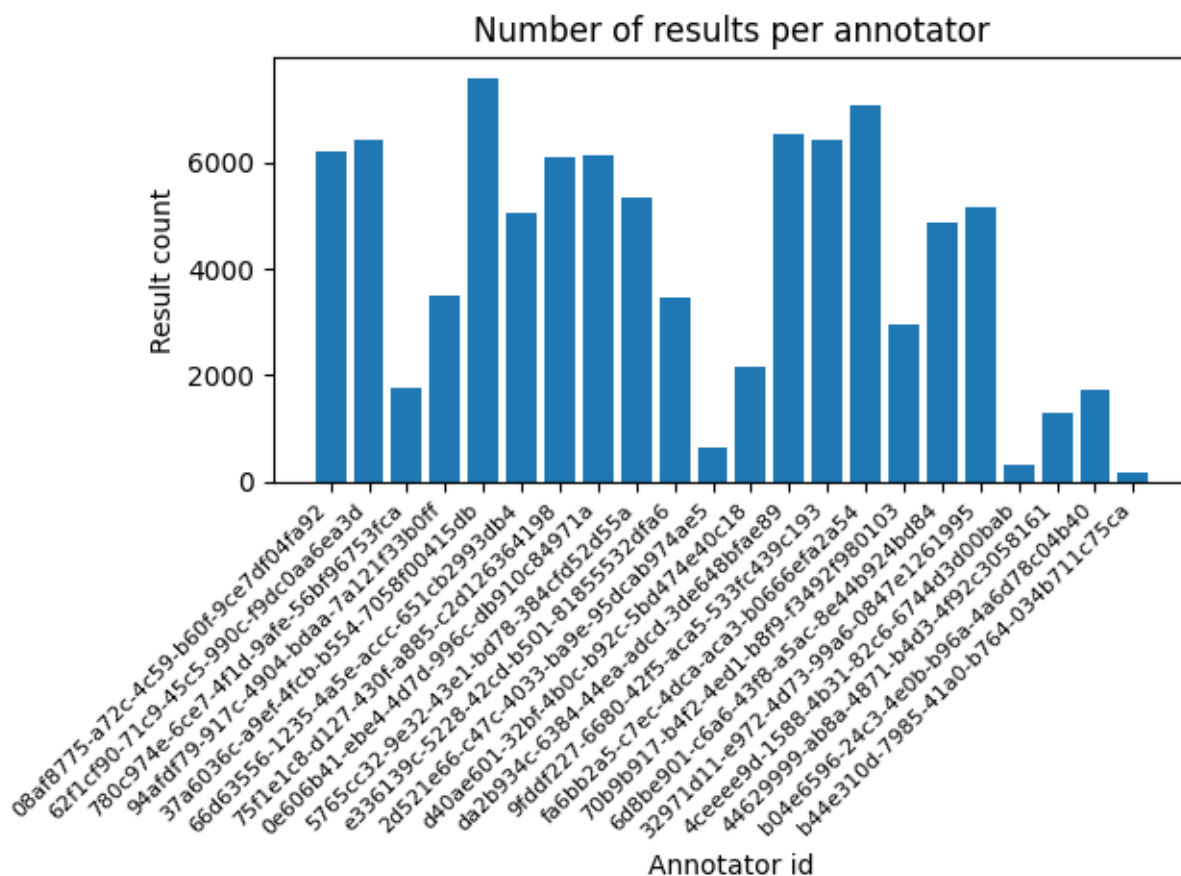


Figure 6: Number of results per annotator

Out of 9087 distinct inputs, 384 of them have a high variance ($v > 0.2$), meaning annotators highly disagree on approximately 4.2% of questions. The table in figure 7 shows the 10 questions with the highest answer variance, sorted in descending order.

input id	V	Answers
2207	0.25	0, 0, 1, 1, 1, 1, 0, 1, 0, 0
5914	0.25	1, 0, 1, 1, 0, 0, 1, 0, 0, 1
0998	0.25	0, 0, 1, 0, 1, 1, 1, 0, 0, 1
2940	0.25	1, 0, 1, 1, 1, 0, 0, 0, 1, 0
8397	0.25	1, 0, 0, 0, 1, 0, 1, 1, 1, 0
3654	0.25	0, 1, 1, 1, 0, 0, 1, 1, 0, 0
5958	0.25	0, 0, 0, 1, 1, 0, 1, 0, 1, 1
0627	0.25	0, 1, 1, 1, 1, 0, 0, 0, 0, 1
0669	0.25	0, 0, 0, 1, 0, 1, 1, 0, 1, 1
6324	0.25	0, 1, 0, 0, 0, 1, 1, 1, 0, 1

Figure 7: Questions with the highest variance with their answers

When treating each task, annotators had the chance to indicate if the data was corrupted, or if they could not solve the task. Out of 90870 tasks, 17 of them could not be solved and 4 of them had corrupt data. All of these tasks answered “no” to the question and had an average annotation time well below average: 838ms for tasks that could not be solved and 461ms for tasks with corrupt data. It is interesting to note that among the 22 annotators present in this experiment, only 10 of them contributed to the 17 tasks mentioned above.

The reference set is balanced, with 49.5% of thruthy values and 50.5% of falsy values (see figure 8).

Repartition of the references values for each input




Figure 8: Repartition of the references values for each input

Using the reference set, we can compare the output of each annotator to the expected answers, and compute an accuracy rate. The vast majority of annotators (90.9%) have an accuracy rate higher than 90% (see figure 9). With an average annotator accuracy rate of 93.5%, the overall performance is good but it would be interesting to identify the best and worst performing annotators.

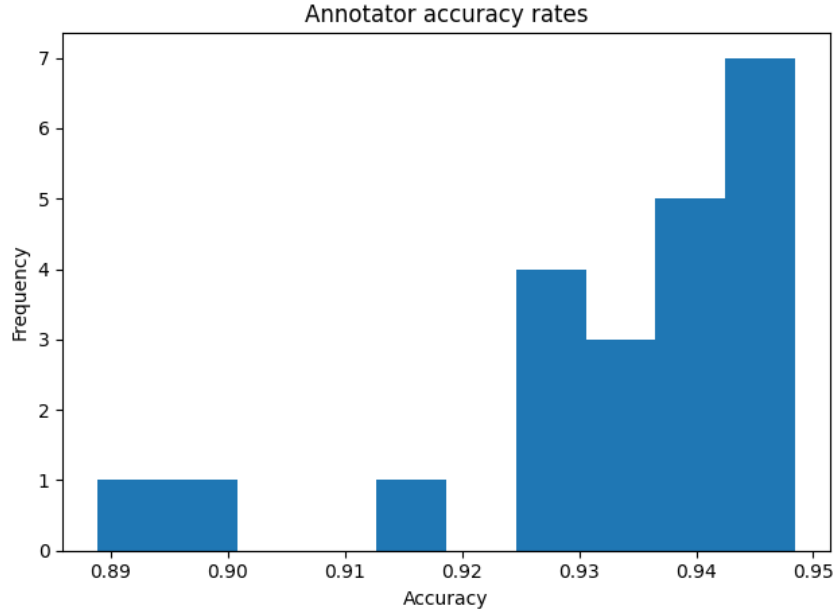


Figure 9: Annotator accuracy rates

Based on the computed accuracy rate of each annotator and their average task duration, we can easily rank them to retrieve the ten worst performing annotators (see figure 10).

Annotator id	Accuracy rate (%)	Average task duration (ms)
4ceeee9d-1588-4b31-82c6-6744d3d00bab	88.89	1253
da2b934c-6384-44ea-adcd-3de648bfae89	89.93	1435
d40ae601-32bf-4b0c-b92c-5bd474e40c18	91.54	1578
5765cc32-9e32-43e1-bd78-384cfd52d55a	92.66	1497
9fddf227-6680-42f5-aca5-533fc439c193	92.84	1114
2d521e66-c47c-4033-ba9e-95dcab974ae5	92.86	1460
08af8775-a72c-4c59-b60f-9ce7df04fa92	92.98	1306
6d8be901-c6a6-43f8-a5ac-8e44b924bd84	93.33	1199
37a6036c-a9ef-4fcb-b554-7058f00415db	93.48	1178
62f1cf90-71c9-45c5-990c-f9dc0aa6ea3d	93.65	1280

Figure 10: Worst performing annotators

Similarly, we can retrieve the ten best performing annotators (see figure 11).

Annotator id	Accuracy rate (%)	Average task duration (ms)
75f1e1c8-d127-430f-a885-c2d126364198	94.84	1365
44629999-ab8a-4871-b4d3-4f92c3058161	94.84	1077
0e606b41-ebe4-4d7d-996c-db910c84971a	94.83	1173
b44e310d-7985-41a0-b764-034b711c75ca	94.71	-1303
fa6bb2a5-c7ec-4dca-aca3-b0666efa2a54	94.69	1155
b04e6596-24c3-4e0b-b96a-4a6d78c04b40	94.67	1595
780c974e-6ce7-4f1d-9afe-56bf96753fca	94.33	879
66d63556-1235-4a5e-acc-651cb2993db4	94.23	1270
94afdf79-917c-4904-bdaa-7a121f33b0ff	94.2	992
70b9b917-b4f2-4ed1-b8f9-f3492f980103	94.0	1239

Figure 11: Best performing annotators

In the two previous figures, the ranks are solely based on the accuracy rate. Ideally, other criteria should also be taken into account: difficulty of the task (depending on the average annotator percentage of correctness), task duration.

We can imagine a ranking system where each annotator is awarded a number of points on completion of a task. Task results with a wrong answer will not be rewarded. The number of points rewarded for each task will be computed based on the following equation:

$$task_points = 10 \left(\frac{avg_task_duration_time}{task_duration_time} \right) (1 + (1 - task_completion_percentage^2))$$

Using this ranking system, the best performing annotators are the one with the highest average points per task (see figure 12).

Annotator id	Average points per task	Rank difference with figure 11
780c974e-6ce7-4f1d-9afe-56bf96753fca	30.87	↑ 6
94afdf79-917c-4904-bdaa-7a121f33b0ff	25.65	↑ 7
44629999-ab8a-4871-b4d3-4f92c3058161	24.55	↓ 1
9fddf227-6680-42f5-aca5-533fc439c193	23.61	↑ 14
62f1cf90-71c9-45c5-990c-f9dc0aa6ea3d	23.46	↑ 8
fa6bb2a5-c7ec-4dca-aca3-b0666efa2a54	23.11	↓ 1
4ceeee9d-1588-4b31-82c6-6744d3d00bab	23.07	↑ 15
0e606b41-ebe4-4d7d-996c-db910c84971a	22.55	↓ 5
70b9b917-b4f2-4ed1-b8f9-f3492f980103	22.44	↑ 1
37a6036c-a9ef-4fcb-b554-7058f00415db	22.29	↑ 4

Figure 12: Best performing annotators based on computed score

In order to better visualize if there is any correlation between annotators accuracy rate and the average task duration, a scatter plot would be ideal. As seen in figure 13, there doesn't seem to be any correlation between the two criterias.

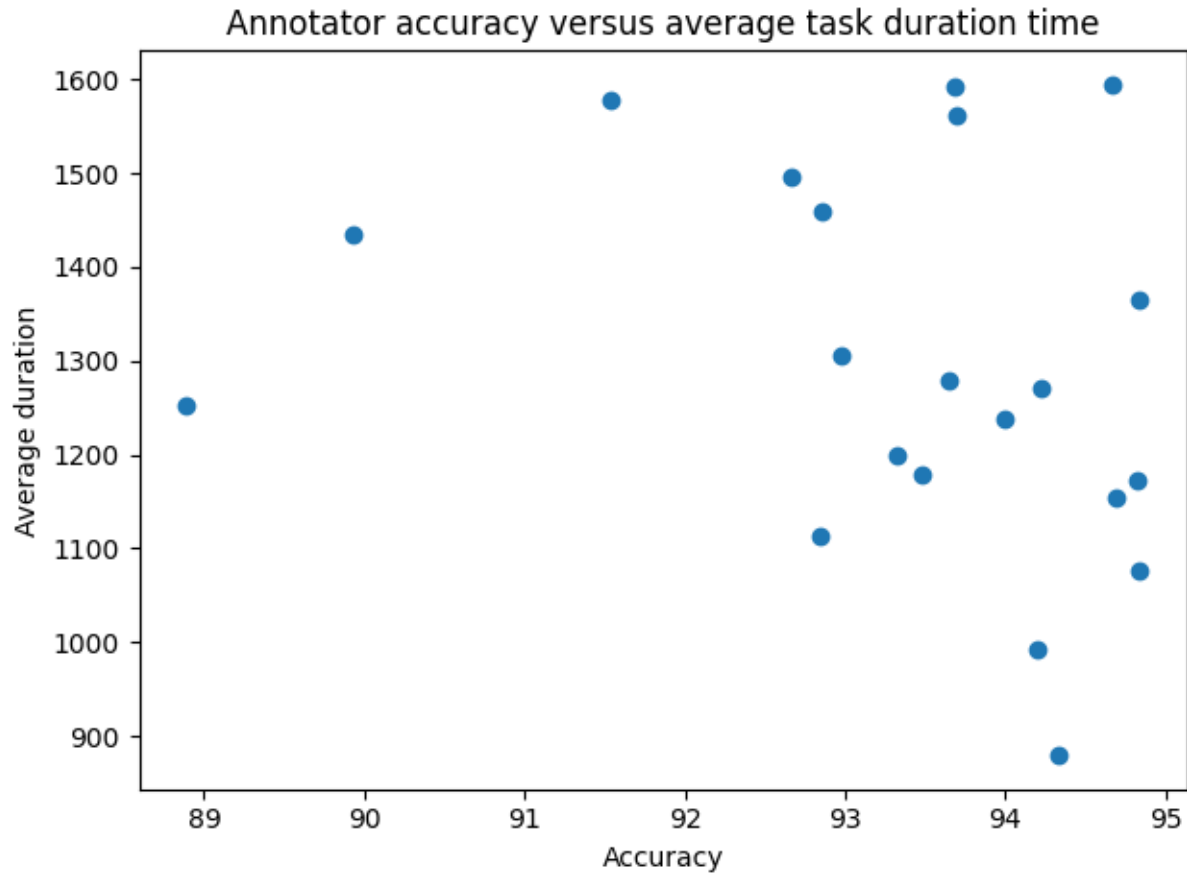


Figure 13: Annotator accuracy rates

5 Conclusions

This experiment allowed us to learn about interesting key performance indicators such as the average processing time of a task, or the performance of annotators based on the reference data.

The moderate size of the JSON file containing the results facilitated the processing and analysis of the data, but the current approach could be problematic with a higher number of results: from a certain threshold, it would no longer be possible to load the complete data set in memory.

To overcome this problem, different options are possible: batch processing, use of a database which is optimized to handle large volumes of data, . The extracted data of the data set being relational, it would be ideal to insert them in a relational database (type Postgresql) in order to be able to execute SQL queries. SQL queries are very fast and flexible and are used across different systems (Postgresql, MySQL, Oracle, Amazon Athena...).

Moreover, the migration to a database would be very easy by using the existing models in the Python code and an ORM. ORMs allow querying and manipulating data in a relational database using an object-oriented paradigm.